

Flattening Irregular Nested Parallelism

Cosmin E. Oancea
`cosmin.oancea@diku.dk`

Department of Computer Science (DIKU)
University of Copenhagen

December 2024 DPP Lecture Slides

Parallel Basic Blocks Recap

Part I: Flattening Nested and Irregular Parallelism

- What is "Flattening"? Recipe for Applying Flattening
- Several Re-Write Rules (inefficient for replicate & iota)
- Jagged (Irregular Multi-Dim) Array Representation
- Revisiting the Rewrites for Replicate & Iota Nested Inside Map
- Revisiting the Solution to Our Example

Part II: Flattening Nested and Irregular Parallelism

- More Flattening Rules
- Flattening by Function Lifting
- Flattening Quicksort
- Flattening Prime-Number (Sieve) Computation

Zip, Unzip, iota, replicate

- $\text{zip} : [n]_{\alpha_1} \rightarrow [n]_{\alpha_2} \rightarrow [n](\alpha_1, \alpha_2)$
- $\text{zip } [a_1, \dots, a_n] [b_1, \dots, b_n] \equiv [(a_1, b_1), \dots, (a_n, b_n)],$
- $\text{unzip} : [n](\alpha_1, \alpha_2) \rightarrow ([n]_{\alpha_1}, [n]_{\alpha_2})$
- $\text{unzip } [(a_1, b_1), \dots, (a_n, b_n)] \equiv ([a_1, \dots, a_n], [b_1, \dots, b_n]),$
- In some sense zip/unzip are syntactic sugar
- $\text{replicate} : (n: \text{int}) \rightarrow \alpha \rightarrow [n]_{\alpha}$
- $\text{replicate } n \ a \equiv [a, a, \dots, a],$
- $\text{iota} : (n: \text{int}) \rightarrow [n]_{\text{int}}$
- $\text{iota } n \equiv [0, 1, \dots, n-1]$

Note: in Haskell zip does not expect same-length arrays;
in Futhark it does!

Map, Reduce, and Scan Types and Semantics

- $[n]\alpha$ denotes the type of an array of n elements of type α .
- $\text{map} : (\alpha \rightarrow \beta) \rightarrow [n]\alpha \rightarrow [n]\beta$
 $\text{map } f [x_1, \dots, x_n] = [f \ x_1, \dots, f \ x_n],$
i.e., $x_i : \alpha, \forall i$, and $f : \alpha \rightarrow \beta$.
- $\text{reduce} : (\alpha \rightarrow \alpha \rightarrow \alpha) \rightarrow \alpha \rightarrow [n]\alpha \rightarrow \alpha$
 $\text{reduce } \odot \ e [x_1, x_2, \dots, x_n] = e \odot x_1 \odot x_2 \odot \dots \odot x_n,$
i.e., $e : \alpha$, $x_i : \alpha, \forall i$, and $\odot : \alpha \rightarrow \alpha \rightarrow \alpha$.
- $\text{scan}^{\text{exc}} : (\alpha \rightarrow \alpha \rightarrow \alpha) \rightarrow \alpha \rightarrow [n]\alpha \rightarrow [n]\alpha$
 $\text{scan}^{\text{exc}} \odot \ e [x_1, \dots, x_n] = [e, e \odot x_1, \dots, e \odot x_1 \odot \dots \odot x_{n-1}]$
i.e., $e : \alpha$, $x_i : \alpha, \forall i$, and $\odot : \alpha \rightarrow \alpha \rightarrow \alpha$.
- $\text{scan}^{\text{inc}} : (\alpha \rightarrow \alpha \rightarrow \alpha) \rightarrow \alpha \rightarrow [n]\alpha \rightarrow [n]\alpha$
 $\text{scan}^{\text{inc}} \odot \ e [x_1, \dots, x_n] = [e \odot x_1, \dots, e \odot x_1 \odot \dots \odot x_n]$
i.e., $e : \alpha$, $x_i : \alpha, \forall i$, and $\odot : \alpha \rightarrow \alpha \rightarrow \alpha$.

Map2, Filter

- $\text{map2} : (\alpha_1 \rightarrow \alpha_2 \rightarrow \beta) \rightarrow [n]\alpha_1 \rightarrow [n]\alpha_2 \rightarrow [n]\beta$
- $\text{map2 } \odot [a_1, \dots, a_n] [b_1, \dots, b_n] \equiv [a_1 \odot b_1, \dots, a_n \odot b_n]$
- $\text{map3 } \dots$
- $\text{filter} : (\alpha \rightarrow \text{Bool}) \rightarrow [n]\alpha \rightarrow [m]\alpha \ (m \leq n)$
- $\text{filter } p [a_1, \dots, a_n] = [a_{k_1}, \dots, a_{k_m}]$ such that $k_1 < k_2 < \dots < k_m$, and denoting $\bar{k} = k_1, \dots, k_m$, we have $(p \ a_j == \text{true}) \ \forall j \in \bar{k}$, **and** $(p \ a_j == \text{false}) \ \forall j \notin \bar{k}$.

Note: in Haskell `map2`, `map3` do not expect same-length arrays; in Futhark they do!

Scatter: A Parallel Write Operator

Scatter **updates in parallel** a base array with a set of values at specified indices:

$\text{scatter} : *[m]\alpha \rightarrow [n]\text{int} \rightarrow [n]\alpha \rightarrow *[m]\alpha$

A (data vector) = [b0, b1, b2, b3]

I (index vector) = [2, 4, 1, -1]

X (input array) = [a0, a1, a2, a3, a4, a5]

$\text{scatter } X \mid A = [a0, b2, b0, a3, b1, a5]$

Scatter: A Parallel Write Operator

Scatter **updates in parallel** a base array with a set of values at specified indices:

scatter : $*[m]_{\alpha} \rightarrow [n]_{\text{int}} \rightarrow [n]_{\alpha} \rightarrow *[m]_{\alpha}$

A (data vector) = [b0, b1, b2, b3]

I (index vector) = [2, 4, 1, -1]

X (input array) = [a0, a1, a2, a3, a4, a5]

scatter X I A = [a0, b2, b0, a3, b1, a5]

scatter has $D(n) = \Theta(1)$ and $W(n) = \Theta(n)$,
i.e., requires n update operations (n is the size of I or A , not of X !).

- 1 Array X is consumed by **scatter**; following uses of X are illegal!
- 2 Similarly, X can alias neither I nor A !

In Futhark, **scatter** checks and ignores the indices that are out of bounds (no update is performed on those). This is useful for padding the iteration space in order to obtain regular parallelism.

Partition2/Filter Implementation

`partition2: ($\alpha \rightarrow \text{Bool}$) \rightarrow $[n]\alpha \rightarrow (\text{i32}, [n]\alpha)$`

In result, the elements satisfying the predicate occur before the others. **Can be implemented by means of map, scan, scatter.**

```
let partition2 't [n] (dummy: t)
  (cond: t -> bool) (X: [n]t) :
  (i64, [n]t) =
```

Assume $X = [5, 4, 2, 3, 7, 8]$, and
cond is T(rue) for even nums.

```
let cs = map cond X
let tfs = map (\ f-> if f then 1
                  else 0) cs

let isT = scan (+) 0 tfs
let i = isT[n-1]

let ffs = map (\ f-> if f then 0
                  else 1) cs
let isF = map (+ i) <| scan (+) 0 ffs
let inds = map (\ (c, iT, iF) ->
                  if c then iT-1
                  else iF-1
                ) (zip3 cs isT isF)

let tmp = replicate n dummy
in (i, scatter tmp inds X)
```


Partition2/Filter Implementation

`partition2: ($\alpha \rightarrow \text{Bool}$) \rightarrow $[n]\alpha \rightarrow (\text{i32}, [n]\alpha)$`

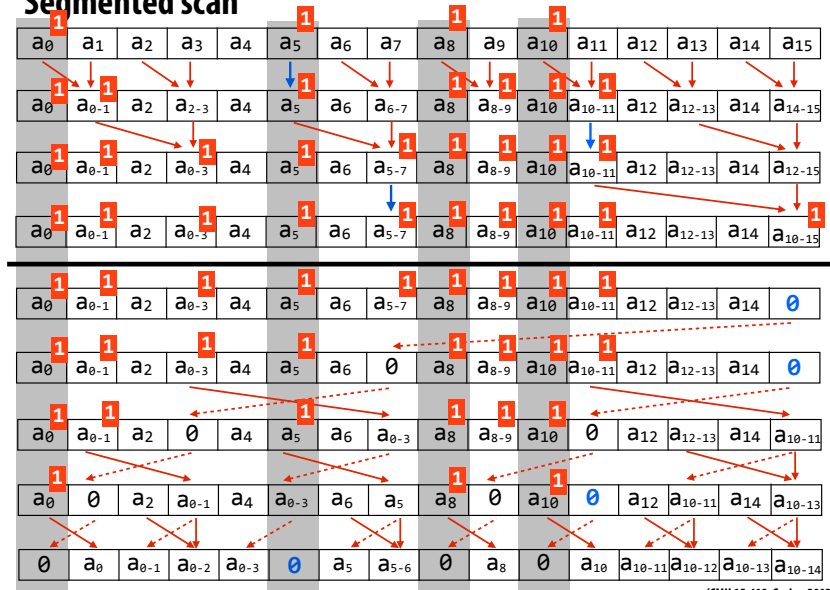
In result, the elements satisfying the predicate occur before the others. **Can be implemented by means of map, scan, scatter.**

```
let partition2 't [n] (dummy: t)
  (cond: t -> bool) (X: [n]t) :
  (i64, [n]t) =
  let cs = map cond X
  let tfs = map (\ f-> if f then 1
                      else 0) cs
  let isT = scan (+) 0 tfs
  let i = isT[n-1]

  let ffs = map (\ f-> if f then 0
                      else 1) cs
  let isF = map (+ i) <| scan (+) 0 ffs
  let inds = map (\(c, iT, iF) ->
                      if c then iT-1
                      else iF-1
                  ) (zip3 cs isT isF)
  let tmp = replicate n dummy
  in (i, scatter tmp inds X)
```

Assume $X = [5, 4, 2, 3, 7, 8]$, and
cond is $T(\text{rue})$ for even nums.
 $n = 6$
 $cs = [F, T, T, F, F, T]$
 $tfs = [0, 1, 1, 0, 0, 1]$
 $isT = [0, 1, 2, 2, 2, 3]$
 $i = 3$
 $ffs = [1, 0, 0, 1, 1, 0]$
 $isF = [4, 4, 4, 5, 6, 6]$
 $inds = [3, 0, 1, 4, 5, 2]$
 $flags = [3, 0, 0, 3, 0, 0]$
 $Result = [4, 2, 8, 5, 3, 7]$

Segmented scan



Segmented Scan Is a Sort of Scan

```
def sgmscan 't [n] (op: t->t->t) (ne: t)
    (flg : [n]bool) (arr : [n]t) : [n]t =
  let flgs_vals =
    zip flg arr |>
    scan ( \ (f1, x1) (f2, x2) ->
      let f = f1 || f2
      in if f2 then (f, x2)
        else (f, op x1 x2)
    ) (false, ne)
  let (_, vals) = unzip flgs_vals
  in vals
```

```
sgmscan (+) 0 [1,0,0,1,0, 0, 0]
              [1,2,3,4,5, 6, 7]
              = = = = =
              [1,3,6,4,9,15,22]
```

```
map (\ row -> scan (+) 0 row)
    [[1,2,3], [4,5, 6, 7]]
    = = = = =
    [[1,3,6], [4,9,15,22]]
```

Correctness Argument:

Segmented Scan Is a Sort of Scan

```
def sgmscan 't [n] (op: t->t->t) (ne: t)
    (flg : [n]bool) (arr : [n]t) : [n]t =
  let flgs_vals =
    zip flg arr |>
    scan ( \ (f1, x1) (f2, x2) ->
      let f = f1 || f2
      in if f2 then (f, x2)
        else (f, op x1 x2)
    ) (false, ne)
  let (_, vals) = unzip flgs_vals
  in vals
```

```
sgmscan (+) 0 [1,0,0,1,0, 0, 0]
               [1,2,3,4,5, 6, 7]
               = = = = =
               [1,3,6,4,9,15,22]
```

```
map (\ row -> scan (+) 0 row)
    [[1,2,3], [4,5, 6, 7]]
    = = = = =
    [[1,3,6], [4,9,15,22]]
```

Correctness Argument:

verify sequential semantics + associative operator \Rightarrow
parallel semantics also holds

Parallel Basic Blocks Recap

Part I: Flattening Nested and Irregular Parallelism

- What is "Flattening"? Recipe for Applying Flattening
- Several Re-Write Rules (inefficient for replicate & iota)
- Jagged (Irregular Multi-Dim) Array Representation
- Revisiting the Rewrites for Replicate & Iota Nested Inside Map
- Revisiting the Solution to Our Example

Part II: Flattening Nested and Irregular Parallelism

- More Flattening Rules
- Flattening by Function Lifting
- Flattening Quicksort
- Flattening Prime-Number (Sieve) Computation

What is "Flattening"?

A code transformation, attributed to Blelloch in the context of the NESL languages, that takes as input a nested parallel program—possibly involving recursion and irregular/jagged arrays—and produces a semantically-equivalent, flat-parallel programs that runs optimally on a PRAM machine.

Meaning: *it is guaranteed to preserve the work and depth of the original nested-parallel program.*^{**}

^{**} As long as scan has $O(1)$ depth and concat has $O(1)$ work and . . .

Flattening Pros and Cons

Pros:

- + clever code transformation
- + important as a programming technique as well (promotes parallel **thinking**)
- + perhaps the only way of mapping a set of challenging problems to capricious architectures such as GPUs (e.g., that do not supports dynamic scheduling of parallelism)

Flattening Pros and Cons

Pros:

- + clever code transformation
- + important as a programming technique as well (promotes parallel **thinking**)
- + perhaps the only way of mapping a set of challenging problems to capricious architectures such as GPUs (e.g., that do not supports dynamic scheduling of parallelism)

Cons:

- does not consider communication/locality and hardware gets more and more heterogeneous
- worse, it tends to destroy the available locality and may explode memory footprint
- useful to cover datasets that fall outside the “common case”

Demonstration at the end of the second Flattening lecture

Flattening: A Bird's Eye View

Incomplete recipe for flattening a nested-parallel program consisting of maps and scan/reduce/scatters at innermost level:

- I. **Normalize the program (think 3-address form).**
The easy way is to also eliminate free variables appearing in the current map if they are variant in an outer, enclosing map.
- II. **Distribute the parallel context (perfect nest of maps) across the enclosed let-binding statements and handle recurrences by function lifting or map-loop interchange.** Systematic application results in a smallish number of code patterns.
- III. **Apply a set of rewrite rules to flatten each pattern**, e.g., treating the cases of a reduce, scan, replicate, iota, scatter, array index which is perfectly nested inside the context.

Flattening: A Bird's Eye View

Incomplete recipe for flattening a nested-parallel program consisting of maps and scan/reduce/scatters at innermost level:

- I. **Normalize the program (think 3-address form).**
The easy way is to also eliminate free variables appearing in the current map if they are variant in an outer, enclosing map.
- II. **Distribute the parallel context (perfect nest of maps) across the enclosed let-binding statements and handle recurrences by function lifting or map-loop interchange.** Systematic application results in a smallish number of code patterns.
- III. **Apply a set of rewrite rules to flatten each pattern**, e.g., treating the cases of a reduce, scan, replicate, iota, scatter, array index which is perfectly nested inside the context.

Differences w.r.t. the PMPH material:

- 1 **also optimize the number of accesses to global memory**
flattening results in memory-bound performance behavior.
- 2 **cover more rewrite rules and more challenging problems**
(e.g., divide-and-conquer recursion)

PMPH Recap: A Simple Demonstration of How to Flatten

Contrived Example:

```
let arr = [1, 2, 3, 4] in  
map (\i -> map (+(i+1)) (iota i)) arr  
-- Result: [[2],[3,4],[4,5,6],[5,6,7,8]]
```

PMPH Recap: A Simple Demonstration of How to Flatten

Contrived Example:

```
let arr = [1, 2, 3, 4] in
map (\i -> map (+(i+1)) (iota i)) arr
-- Result: [[2],[3,4],[4,5,6],[5,6,7,8]]
```

I. Normalize the code:

```
map (\i -> let ip1 = i+1 in
            let iot = (iota i) in
            let ip1r = (replicate i ip1)
            in map2 (+) ip1r iot          ) arr
```

PMPH Recap: A Simple Demonstration of How to Flatten

Contrived Example:

```
let arr = [1, 2, 3, 4] in
map (\i -> map (+(i+1)) (iota i)) arr
-- Result: [[2],[3,4],[4,5,6],[5,6,7,8]]
```

I. Normalize the code:

```
map (\i -> let ip1 = i+1 in
            let iot = (iota i) in
            let ip1r = (replicate i ip1)
            in map2 (+) ip1r iot                ) arr
```

II. Distribute the map across every statement in the body

and adjust the inputs accordingly (\mathcal{F} denotes the transformation)

```
 $\mathcal{F}(\text{map } (\lambda i \rightarrow \text{map } (+(i+1)) \text{ (iota } i)) \text{ arr}) \equiv$ 
1. let ip1s = map (\i -> i+1) arr in -- [2, 3, 4, 5]
2. let iots =  $\mathcal{F}(\text{map } (\lambda i \rightarrow \text{iota } i)) \text{ arr}$  in
3. let ip1rs=  $\mathcal{F}(\text{map2 } (\lambda i \text{ ip1} \rightarrow \text{replicate } i \text{ ip1})) \text{ arr ip1s}$ 
4. in  $\mathcal{F}(\text{map2 } (\lambda \text{ ip1r iot} \rightarrow \text{map2 } (+) \text{ ip1r iot}) \text{ ip1rs iots})$ 
```

For simplicity we assume `arr` contains strictly-positive integers.

PMPH Recap: A Simple Demonstration of How to Flatten

According to inefficient rule “iota nested inside a map”

(assuming `arr = [1,2,3,4]`):

```
2. let iots =  $\mathcal{F}$ (map (\i -> iota i) arr)
```

≡

```
inds = scanexc (+) 0 arr           -- [0,1,3,6]
size = (last inds) + (last arr)    -- 6 + 4 = 10
flag = scatter (replicate size 0)  -- [1, 2, 0, 3, 0, 0, 4, 0, 0, 0]
      inds arr
tmp  = replicate size 1
iots = sgmScanexc (+) 0 flag tmp    -- [0, 0, 1, 0, 1, 2, 0, 1, 2, 3]
```

PMPH Recap: A Simple Demonstration of How to Flatten

According to inefficient rule “replicate nested inside a map”

(assuming `arr = [1,2,3,4]`):

```
3. let ip1rs=  $\mathcal{F}$ (map2 (\ i ip1 -> replicate i ip1) arr ip1s)
≡
vals = scatter (replicate size 0) inds ip1s -- [2,3,0,4,0,0,5,0,0,0]
ip1rs= sgmScaninc (+) 0 flag vals -- [2,3,3,4,4,4,5,5,5,5]
```

According to rule “map nested inside a map”

```
 $\mathcal{F}$ (map2 (\ ip1r iot -> map2 (+) ip1r iot) ip1rs iots)
≡
4. result = map2 (+) ip1rs iots
-- [2, 3, 3, 4, 4, 4, 5, 5, 5, 5]
-- [0, 0, 1, 0, 1, 2, 0, 1, 2, 3]
-- + + + + + + + + + +
-----
-- [2, 3, 4, 4, 5, 6, 5, 6, 7, 8] values
%-- [1, 2, 0, 3, 0, 0, 4, 0, 0, 0] flags
```

At each step we also reason about the shape of the resulting array.
The shape of the 2D jagged arrays `iots`, `ip1rs`, `result` **is** `arr`.

Parallel Basic Blocks Recap

Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening

Several Re-Write Rules (inefficient for replicate & iota)

Jagged (Irregular Multi-Dim) Array Representation

Revisiting the Rewrites for Replicate & Iota Nested Inside Map

Revisiting the Solution to Our Example

Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules

Flattening by Function Lifting

Flattening Quicksort

Flattening Prime-Number (Sieve) Computation

Nested vs Flattened Parallelism: Scan inside a Map

(1) Scan nested inside a map:

```
res = map (\row->scaninc (+) 0 row) [[1,3], [2,4,6]]
```

≡

```
res = [ scaninc (+) 0 [1,3],      scaninc (+) 0 [2,4,6] ]
```

≡

```
res = [ [ 1, 4],                  [2, 6, 12] ]
```

Nested vs Flattened Parallelism: Scan inside a Map

(1) Scan nested inside a map:

```
res = map (\row->scaninc (+) 0 row) [[1,3], [2,4,6]]  
≡  
res = [ scaninc (+) 0 [1,3],      scaninc (+) 0 [2,4,6] ]  
≡  
res = [ [ 1, 4],                  [2, 6, 12] ]
```

becomes a segmented scan, which requires a flag array as arg:

```
sgmScaninc (+) 0 [1, 0, 1, 0, 0] [1, 3, 2, 4, 6] ≡ [ 1, 4, 2, 6, 12 ]
```

Flattening a scan directly nested inside a map:

- $S_{arr}^1, F_{arr}, D_{arr}$ denote the shape, flag & flat data of input `arr`.
- The flat-data result is obtained by a segmented scan.
- The shape of the result array is the same as the input array.

$\mathcal{F}(\text{res} = \text{map } (\backslash \text{row} \rightarrow \text{scan } (\odot) 0_{\odot} \text{ row}) \text{ arr}) \Rightarrow$

$S_{res}^1 = S_{arr}^1$

$D_{res} = \text{sgmScan } (\odot) 0_{\odot} F_{arr} D_{arr}$

Nested vs Flattened Parallelism: Map inside a Map

(2) Map nested inside a map:

```
res = map (\row->map f row) [[1,3], [2,4,6]]
```

≡

```
res = [ map f [1, 3],      map f [2, 4, 6] ]
```

≡

```
res = [ [f(1),f(3)], [f(2),f(4),f(6)] ]
```

Nested vs Flattened Parallelism: Map inside a Map

(2) Map nested inside a map:

```
res = map (\row->map f row) [[1,3], [2,4,6]]  
≡  
res = [ map f [1, 3],      map f [2, 4, 6] ]  
≡  
res = [ [f(1),f(3)], [f(2),f(4),f(6)] ]
```

Flattening a map directly nested inside a map:

- the flat-data array is obtained by a map on the flat input;
- the shape of the result array is the same as the input array.

$\mathcal{F}(\text{res} = \text{map } (\lambda \text{row} \rightarrow \text{map } f \text{ row}) \text{ arr}) \Rightarrow$

$$S_{\text{res}}^1 = S_{\text{arr}}^1$$

$$D_{\text{res}} = \text{map } f \ D_{\text{arr}}$$

Nested vs Flattened Parallelism: Replicate inside Map

(3) Replicate nested inside a map:

```
res = map2 (\ n m -> replicate n m) [1,0,3,2] [7,3,8,9] ≡  
res = [ replicate 1 7, replicate 0 3, replicate 3 8, replicate 2 9 ] ≡  
res = [ [7], [], [8,8,8], [9,9] ]
```

Nested vs Flattened Parallelism: Replicate inside Map

(3) Replicate nested inside a map:

```
res = map2 (\ n m -> replicate n m) [1,0,3,2] [7,3,8,9] ≡  
res = [ replicate 1 7, replicate 0 3, replicate 3 8, replicate 2 9 ] ≡  
res = [ [7], [], [8,8,8], [9,9] ]
```

```
res = map2(\n m-> replicate n m) ns ms  
becomes a scan-scatter composition:
```

Nested vs Flattened Parallelism: Replicate inside Map

(3) Replicate nested inside a map:

```
res = map2 (\ n m -> replicate n m) [1,0,3,2] [7,3,8,9] ≡  
res = [ replicate 1 7, replicate 0 3, replicate 3 8, replicate 2 9 ] ≡  
res = [ [7], [], [8,8,8], [9,9] ]
```

```
res = map2(\n m-> replicate n m) ns ms
```

becomes a scan-scatter composition:

1. the shape of the result array is ns
- 2-3. builds the indices at which segment start (-1 for null shape)
4. get the size of the flat array (summing ns)
- 5-6. write the ms and ns values at the start of their segments
7. propagate the ms values throughout their segments.

```
 $\mathcal{F}$ (res = map2 (\n m -> replicate n m) ns ms)  $\Rightarrow$ 
```

1. $S_{res}^1 = ns$
2. inds = scan^{exc} (+) 0 ns
3. |> map2 (\n i->if n>0 then i else -1) ns
4. size = (last inds) + (last ns)
5. vls = scatter (replicate size 0) inds ms
6. F_{arr} = scatter (replicate size 0) inds ns
6. D_{res} = sgmScan^{inc} (+) 0 F_{res} vls

```
-- ms = [7,3,8,9]  
-- ns = [1,0,3,2]  
-- [0,1,1,4]  
-- inds = [0,-1,1,4]  
-- 4 + 2 = 6  
-- [7, 8, 0, 0, 9, 0]  
-- [1, 3, 0, 0, 2, 0]  
-- [7, 8, 8, 8, 9, 9]
```

Nested vs Flattened Parallelism: Iota inside Map

(4) Iota nested inside a map $((\text{iota } n) \equiv [0, \dots, n-1])$:

```
res = map (\i -> iota i) [1,3,2]  $\equiv$ 
```

```
res = [iota 1, iota 3, iota 2]  $\equiv$  [ [0], [0,1,2], [0,1] ]
```


Nested vs Flattened Parallelism: Iota inside Map

(4) Iota nested inside a map ($\text{iota } n \equiv [0, \dots, n-1]$):

```
res = map (\i -> iota i) [1,3,2]  $\equiv$   
res = [iota 1, iota 3, iota 2]  $\equiv$  [[0], [0,1,2], [0,1]]
```

boils down to a segmented scan applied to an array of ones:

1. by definition of iota , ns contains the size of each subarray, hence the shape of the result is ns ;
- 2-3. the flag-array of the result, F_{res} , is constructed from ns ; (we will introduce function mkFlagArray a bit later).
4. the result is obtained by an exclusive segmented scan operation applied to an array of ones.

```
 $\mathcal{F}(\text{res} = \text{map } (\backslash n \rightarrow \text{iota } n) \text{ ns}) \Rightarrow$ 
```

1. $S_{\text{res}}^1 = \text{ns}$ -- $\text{ns} = [1, 3, 2]$
2. $\text{trues} = \text{replicate } (\text{length } \text{ns}) \text{ true}$
3. $(-, F_{\text{res}}) = \text{mkFlagArray } \text{ns } 0 \text{ trues}$ -- $F_{\text{res}} = [1, 1, 0, 0, 1, 0]$
4. $D_{\text{res}} = \text{sgmScan}^{\text{exc}} (+) 0 F_{\text{res}} (\text{replicate } \text{flen}_{\text{res}} 1)$ -- $[0, 0, 1, 2, 0, 1]$

Note 1: $\text{iota } n \equiv \text{scan}^{\text{exc}} (+) 0 (\text{replicate } n 1)$.

Note 2: 1 and 0 denote true and false; flen_{res} is the sum of ns .

Parallel Basic Blocks Recap

Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening
Several Re-Write Rules (inefficient for replicate & iota)

Jagged (Irregular Multi-Dim) Array Representation

Revisiting the Rewrites for Replicate & Iota Nested Inside Map
Revisiting the Solution to Our Example

Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules

Flattening by Function Lifting

Flattening Quicksort

Flattening Prime-Number (Sieve) Computation

Shape-Based Representation

- Two dimensional arrays:

```
arr = [ [1,2,3], [4], [], [5,6] ]
```

\Rightarrow

$S_{arr}^0 = [4]$

$S_{arr}^1 = [3, 1, 0, 2]$

$D_{arr} = [1, 2, 3, 4, 5, 6]$

Shape-Based Representation

- Two dimensional arrays:

```
arr = [ [1,2,3], [4], [], [5,6] ]
```

⇒

```
Sarr0 = [4]
```

```
Sarr1 = [3, 1, 0, 2]
```

```
Darr = [1, 2, 3, 4, 5, 6]
```

- Three dimensional arrays:

```
arr = [ [], [ [1,2,3], [4], [], [5,6] ], [ [7], [], [8,9,10] ] ]
```

⇒

Shape-Based Representation

- Two dimensional arrays:

```
arr = [ [1,2,3], [4], [], [5,6] ]
```

⇒

```
Sarr0 = [4]
```

```
Sarr1 = [3, 1, 0, 2]
```

```
Darr = [1, 2, 3, 4, 5, 6]
```

- Three dimensional arrays:

```
arr = [ [], [ [1,2,3], [4], [], [5,6] ], [ [7], [], [8,9,10] ] ]
```

⇒

```
Sarr0 = [3]
```

```
Sarr1 = [0, 4, 3]
```

```
Sarr2 = [3, 1, 0, 2, 1, 0, 3]
```

```
flenarr = 10
```

```
Darr = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

Assume a n-dimensional array; The following invariant holds:

length $S_{arr}^i = \text{reduce } (+) \ 0 \ S_{arr}^{i-1}, \forall 1 \leq i < n$

length $D_{arr} = \text{reduce } (+) \ 0 \ S_{arr}^{n-1}$

Flat Representation: Auxiliary Structures

```
arr = [ [], [ [1,2,3], [4], [], [5,6] ], [ [7], [], [8,9,10] ]
```

⇒

$S_{arr}^0 = [3]$

$S_{arr}^1 = [0, 4, 3]$

$S_{arr}^2 = [3, 1, 0, 2, 1, 0, 3]$

$D_{arr} = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]$

- **Offset Indices (B):** segment-start offset in the flat data:

$B_{arr}^1 = [0, 0, 6]$

$B_{arr}^2 = [0, 3, 4, 4, 6, 7, 7]$

- **Flag Array (F):** start of a segment indicated by a true value (could also use $\neq 0$ integrals), e.g., used for segmented scans:

$F_{arr}^1 = [1, 0, 0, 0, 0, 0, 1, 0, 0, 0]$

$F_{arr}^2 = [1, 0, 0, 1, 1, 0, 1, 1, 0, 0]$

- **Segment and Inner indices (II):**

$II_{arr}^1 = [1, 1, 1, 1, 1, 1, 2, 2, 2, 2]$

$II_{arr}^2 = [0, 0, 0, 1, 3, 3, 0, 2, 2, 2]$

$II_{arr}^3 = [0, 1, 2, 0, 0, 1, 0, 0, 1, 2]$

Auxiliary Structures: Intuitive Motivation

Auxiliary structures are useful to optimize the replication of values.

Nested-Execution Example:

```
let xss = [ [1,2,3], [], [5,7] ]  
let ys  = [ 4, 2, 1 ]  
let rss = map2 (\ xs y -> map (+y) xs ) xss ys
```

⇒

```
rss = [ map (+4) [1,2,3], map (+2) [], map (+1) [5,7] ]  
rss = [ [5,6,7], [], [6,8] ]
```

Traditional flattening would replicate the values of y:

```
let (S1yss, Dyss) =  $\mathcal{F}$ (map2 (\ xs y -> replicate (length xs) y) xss ys)  
let Drss = map2 (\ x y -> x + y) Dxss Dyss
```

⇒

```
Dxss = [1, 2, 3, 5, 7]  
      + + + + +  
Dyss = [4, 4, 4, 1, 1]  
      = = = = =  
Drss = [5, 6, 7, 6, 8]
```

Auxiliary Structures: Intuitive Motivation

Auxiliary structures are useful to optimize the replication of values.

Nested-Execution Example:

```
let xss = [ [1,2,3], [], [5,7] ]  
let ys  = [ 4, 2, 1 ]  
let rss = map2 (\ xs y -> map (+y) xs ) xss ys  
⇒  
rss = [ map (+4) [1,2,3], map (+2) [], map (+1) [5,7] ]  
rss = [ [5,6,7], [], [6,8] ]
```

Using the auxiliary structures we indirectly access other arrays:

```
let Drss = map2 (\ x sgmind -> x + ys[sgmind]) Dxss II1rss  
⇒  
S1rss = [3, 0, 2]  
II1rss = [0, 0, 0, 2, 2]  
Dxss = [1, 2, 3, 5, 7]  
Drss = [1+4, 2+4, 3+4, 5+1, 7+1] = [5, 6, 7, 6, 8]
```

But what have we gained? Creating II_{rss}^1 is as expensive as xss (or better said the expanded yss from the other slide) ...

Auxiliary Structures: Intuitive Motivation

Auxiliary structures are useful to optimize replication:

- they depend only on the shape of the result (created once)
- can indirectly access several lower-dimensional arrays, sharing parallel dimensions!

Nested-Execution Example:

```
let xss = [ [1,2,3], [], [5,7] ]
let ys  = [ 4, 2, 1 ]
let zs  = [ 1, 2, 3 ]
let rss = map3 (\ xs y z -> map (\ x -> x*y + z ) xs ) xss ys zs
⇒
rss = [ [5,9,13], [], [8,10] ]
```

Using the auxiliary structures we indirectly access other arrays:

```
let Drss = map2 (\ y sgmind -> x*ys[sgmind] + zs[sgmind]) Dxss II1rss
⇒
II1rss = [0, 0, 0, 2, 2]
Dxss = [1, 2, 3, 5, 7]
Drss = [1*4+1, 2*4+1, 3*4+1, 5*1+3, 7*1+3] = [5, 9, 13, 8, 10]
```

**We build II_{rss}^1 once and reuse it twice. Also improves locality:
ys and zs are much smaller than xss, hence reused from L1/2\$.**

Auxiliary Structures: Intuitive Motivation

Nested-Execution Example:

```
let xss = [ [1,3], [2] ]  
let yss = [ [2], [4,5] ]  
let rss = map2 (\xs ys -> map (\x -> map (+x) ys ) xs ) xss yss  
⇒  
rss = [ [[3],[5]], [[6,7]] ]
```

Using the auxiliary structures we indirectly access other arrays:

```
let Drss = map3(\ s1 s2 s3 -> let ind_x = B1xss[s1] + s2  
                             let ind_y = B1yss[s1] + s3  
                             in X[ind_x] + Y[ind_y]  
                             ) ||1rss ||2rss ||3rss
```

⇒

```
B1xss = [0, 2]  
B1yss = [0, 1]  
||1rss = [0, 0, 1, 1]  
||2rss = [0, 1, 0, 0]  
||3rss = [0, 0, 0, 1]  
Drss = [ Dxss[0+0]+Dyss[0+0], Dxss[0+1]+Dyss[0+0]  
         , Dxss[2+0]+Dyss[1+0], Dxss[2+0]+Dyss[1+1] ]  
Drss = [ 1+2, 3+2, 2+4, 2+5 ] = [ 3, 5, 6, 7 ]
```

Constructing the Offset Indices (B)

```
arr = [ [], [ [1,2,3], [4], [], [5,6] ], [ [7], [], [8,9,10] ]
```

⇒

```
S0arr = [3]
```

```
S1arr = [0, 4, 3]
```

```
S2arr = [3, 1, 0, 2, 1, 0, 3]
```

```
Darr = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

Offset Indices (B): segment-start offset in the flat data:

```
B1arr = [0, 0, 6]
```

```
B2arr = [0, 3, 4, 4, 6, 7, 7]
```

How to construct Offset Indices (B)?

Constructing the Offset Indices (B)

```
arr = [ [], [ [1,2,3], [4], [], [5,6] ], [ [7], [], [8,9,10] ]
```

⇒

```
S0arr = [3]
```

```
S1arr = [0, 4, 3]
```

```
S2arr = [3, 1, 0, 2, 1, 0, 3]
```

```
Darr = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

Offset Indices (B): segment-start offset in the flat data:

```
B1arr = [0, 0, 6]
```

```
B2arr = [0, 3, 4, 4, 6, 7, 7]
```

How to construct Offset Indices (B)?

By exclusive scanning the corresponding shape and reindexing!

```
B2arr = scanexc (+) 0 S2arr      -- [0, 3, 4, 4, 6, 7, 7]
```

```
B1arr = scanexc (+) 0 S1arr      -- [0, 0, 4]  
      |> map (\i -> B2arr[i])    -- [0, 0, 6]
```

Constructing the Flag Array

From now on, we discuss only TWO-dimensional irregular arrays!

```
def mkFlagArray 't [m]
  (aoa_shp: [m]u32) (zero: t)      -- aoa_shp=[0,3,1,0,4,2,0]
  (aoa_val: [m]t) : ([m]u32, []t) = -- aoa_val=[1,1,1,1,1,1,1]
  let shp_rot = map (\i->if i==0 then 0 -- shp_rot=[0,0,3,1,0,4,2]
                     else aoa_shp[i-1]
                     ) (iota m)
  let shp_scn = scan (+) 0 shp_rot -- shp_scn=[0,0,3,4,4,8,10]
  let aoa_len = if m == 0 then 0i64 -- aoa_len = 10
                 else i64.u32 <|
                   shp_scn[m-1]+aoa_shp[m-1]
  let shp_ind = map2 (\shp ind -> -- shp_ind=
                     if shp==0 then -1i64 -- [-1,0,3,-1,4,8,-1]
                     else i64.u32 ind -- scatter
                     ) aoa_shp shp_scn -- [0,0,0,0,0,0,0,0,0,0]
  let r = scatter (replicate aoa_len zero) -- [-1,0,3,-1,4,8,-1]
              shp_ind aoa_val -- [ 1,1,1, 1,1,1, 1]
  in (shp_scn, r) -- r=[1,0,0,1,1,0,0,0,1,0]
```

Versatile: computes B^1 and F^1 of a 2D jagged array of shape `aoa_shp`, with the start-segment values taken from `aoa_val`.

Unless you have a good reason, F should be a bool array (to reduce memory traffic).

Constructing the Segment and Inner Indices

From now on, we discuss only TWO-dimensional irregular arrays!

```
arr = [ [1,2,3], [4], [], [5,6], [7], [], [8,9,10] ]
```

⇒

```
Sarr0 = [7]
```

```
Sarr1 = [3, 1, 0, 2, 1, 0, 3]
```

```
flenarr = reduce (+) 0 Sarr1 = 10
```

```
Darr = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

Segment and Inner indices (II):

```
IIarr1 = [0, 0, 0, 1, 3, 3, 4, 6, 6, 6]
```

```
IIarr2 = [0, 1, 2, 0, 0, 1, 0, 0, 1, 2]
```

Constructing Segment and Inner indices (II):

```
(Barr1, Farr) = mkFlagArray Sarr1 0 (iota (length Sarr1))  
-- ([0, 3, 4, 4, 6, 7, 7], [0, 0, 0, 1, 3, 0, 4, 6, 0, 0])
```

```
IIarr1 = ???
```

```
IIarr2 = ???
```

Constructing the Segment and Inner Indices

I need to get this:

$II_{arr}^1 = [0, 0, 0, 1, 3, 3, 4, 6, 6, 6]$

from this:

```
(-, Farr) = mkFlagArray Sarr1 0 (iota (length Sarr1))  
-- [0, 0, 0, 1, 3, 0, 4, 6, 0, 0]
```

How?

Constructing the Segment and Inner Indices

I need to get this:

$$II_{arr}^2 = [0, 1, 2, 0, 0, 1, 0, 0, 1, 2]$$

from these:

$$B_{arr} = [0, 3, 4, 4, 6, 7, 7]$$

$$II_{arr}^1 = [0, 0, 0, 1, 3, 3, 4, 6, 6, 6]$$

$$iota\ 10 = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]$$

II_{arr}^1 and II_{arr}^2 have the same length as flat `arr`, in our case 10.

How?

We can also construct it by binary searching B_{arr} or by means of a segmented scan.

Constructing the Segment and Inner Indices

From now on, we discuss only TWO-dimensional irregular arrays!

```
arr = [ [1,2,3], [4], [], [5,6], [7], [], [8,9,10] ]
```

⇒

```
S0arr = [7]
```

```
S1arr = [3, 1, 0, 2, 1, 0, 3]
```

```
flenarr = reduce (+) 0 S1arr = 10
```

```
Darr = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

Segment and Inner indices (II):

```
II1arr = [0, 0, 0, 1, 3, 3, 4, 6, 6, 6]
```

```
II2arr = [0, 1, 2, 0, 0, 1, 0, 0, 1, 2]
```

Constructing Segment and Inner indices (II):

```
(B1arr, F'arr) = mkFlagArray S1arr 0 (iota (length S1arr))  
-- ([0, 3, 4, 4, 6, 7, 7], [0, 0, 0, 1, 3, 0, 4, 6, 0, 0])
```

```
Farr = map bool.u32 F'arr
```

```
II1arr = sgmScaninc (+) 0 Farr F'arr
```

```
II2arr = map2 (\ i sgm -> i - B1arr[sgm] ) (iota flenarr) II1arr  
-- ^ this fuses better & performs less memory traffic than the below:
```

```
II2arr = sgmScaninc (+) 0 Farr (replicate flen 1) |> map (-1)
```

B^{inc} and II^1 Are the Important Ones

Because you can deduce the other arrays by means of simple maps, that fuse better and generate less traffic.

```
arr = [ [1,2,3], [4], [], [5,6], [7], [], [8,9,10] ]
```

⇒

```
 $S_{arr}^0$  = [7]
```

```
 $S_{arr}^1$  = [3, 1, 0, 2, 1, 0, 3]
```

```
flenarr = reduce (+) 0  $S_{arr}^1$  = 10
```

```
 $D_{arr}$  = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

If we know the Segment Offsets (B_{arr}^{inc}) and Indices (II_{arr}^1):

B_{arr}^{inc} = [3, 4, 4, 6, 7, 7, 10] -- *inclusive scan of S_{arr}^1 and*

II_{arr}^1 = [0, 0, 0, 1, 3, 3, 4, 6, 6, 6]

We can efficiently compute the S_{arr}^1 and F_{arr} arrays (also II_{arr}^2) by:

```
 $S_{arr}^1$  = iota (length  $B_{arr}^{inc}$ )
```

```
> map (\i -> if i == 0 then  $B_{arr}^{inc}[i]$  else  $B_{arr}^{inc}[i] - B_{arr}^{inc}[i-1]$ )
```

```
 $F_{arr}$  = iota flenarr -- flenarr is the length of  $II_{arr}^1$ 
```

```
> map (\i -> if i == 0 then true else  $II_{arr}^1[i] \neq II_{arr}^1[i-1]$ )
```

Note: B_{arr}^{inc} different than B_{arr}^1 : inclusive vs exclusive scan of shape!

Parallel Basic Blocks Recap

Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening

Several Re-Write Rules (inefficient for replicate & iota)

Jagged (Irregular Multi-Dim) Array Representation

Revisiting the Rewrites for Replicate & Iota Nested Inside Map

Revisiting the Solution to Our Example

Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules

Flattening by Function Lifting

Flattening Quicksort

Flattening Prime-Number (Sieve) Computation

Revisiting Replicate inside Map

(3) Replicate nested inside a map:

```
res = map2 (\ n m -> replicate n m) [1,0,3,2] [7,3,8,9] ≡  
res = [ replicate 1 7, replicate 0 3, replicate 3 8, replicate 2 9 ] ≡  
res = [ [7], [], [8,8,8], [9,9] ]
```

```
res = map2(\n m-> replicate n m) ns ms
```

becomes a very simple gather operation:

1. the shape of the result array is ns
2. build II^1 of a jagged array of shape ns

Revisiting Replicate inside Map

(3) Replicate nested inside a map:

```
res = map2 (\ n m -> replicate n m) [1,0,3,2] [7,3,8,9] ≡  
res = [ replicate 1 7, replicate 0 3, replicate 3 8, replicate 2 9 ] ≡  
res = [ [7], [], [8,8,8], [9,9] ]
```

```
res = map2(\n m-> replicate n m) ns ms
```

becomes a very simple gather operation:

1. the shape of the result array is ns
2. build II^1 of a jagged array of shape ns
3. gather the corresponding values from ms by indexing through II^1

```
 $\mathcal{F}$ (res = map2 (\n m -> replicate n m) ns ms)  $\Rightarrow$       -- ms = [7,3,8,9]  
1.  $S_{res}^1 = ns$                                            -- ns = [1,0,3,2]  
2.  $II_{res}^1 = \dots$  -- construct  $II^1$  for a jagged array of shape ns  
                                   -- [0, 2,2,2, 3,3]  
3.  $D_{res} = \text{map } (\text{sgm} \rightarrow ms[\text{sgm}]) II_{res}^1$           -- [7, 8,8,8, 9,9]
```

Nested vs Flattened Parallelism: Iota inside Map

(4) Iota nested inside a map ($\text{iota } n \equiv [0, \dots, n-1]$):

```
res = map (\i -> iota i) [1,3,2]  $\equiv$   
res = [iota 1, iota 3, iota 2]  $\equiv$  [ [0], [0,1,2], [0,1] ]
```

```
res = map (\n -> iota n) ns
```

The result is exactly the II^2 array of a jagged array of shape ns

$\mathcal{F}(\text{res} = \text{map } (\lambda n \rightarrow \text{iota } n) \text{ ns}) \Rightarrow$

1. $S_{\text{res}}^1 = \text{ns}$ $-- \text{ ns} = [1, 3, 2]$
2. $\text{II}_{\text{res}}^2 = \dots$ $-- \text{construct } \text{II}^2 \text{ of a jagged array of shape ns}$
4. $D_{\text{res}} = \text{II}_{\text{res}}^2$

Parallel Basic Blocks Recap

Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening
Several Re-Write Rules (inefficient for replicate & iota)
Jagged (Irregular Multi-Dim) Array Representation
Revisiting the Rewrites for Replicate & Iota Nested Inside Map
Revisiting the Solution to Our Example

Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules
Flattening by Function Lifting
Flattening Quicksort
Flattening Prime-Number (Sieve) Computation

Revisiting Our Demonstration of How to Flatten

Contrived Example:

```
let arr = [1, 2, 3, 4] in
map (\i -> map (+(i+1)) (iota i)) arr
-- Result: [[2],[3,4],[4,5,6],[5,6,7,8]]
```

I. Normalize the code:

```
map (\i -> let ip1 = i+1 in
           let iot = (iota i) in
           let ip1r = (replicate i ip1)
           in map2 (+) ip1r iot                ) arr
```

II. Distribute the map across every statement in the body

and adjust the inputs accordingly (\mathcal{F} denotes the transformation)

```
 $\mathcal{F}(\text{map } (\lambda i \rightarrow \text{map } (+(i+1)) \text{ (iota } i)) \text{ arr}) \equiv$ 
1. let ip1s = map (\i -> i+1) arr in -- [2, 3, 4, 5]
2. let iots =  $\mathcal{F}(\text{map } (\lambda i \rightarrow \text{iota } i)) \text{ arr}$  in
3. let ip1rs=  $\mathcal{F}(\text{map2 } (\lambda i \text{ ip1} \rightarrow \text{replicate } i \text{ ip1})) \text{ arr ip1s}$ 
4. in  $\mathcal{F}(\text{map2 } (\lambda \text{ ip1r iot} \rightarrow \text{map2 } (+) \text{ ip1r iot}) \text{ ip1rs iots})$ 
```

We do **not** assume that arr contains strictly-positive integers.

Revisiting Our Example: Think Like a Compiler

```
 $\mathcal{F}(\text{map } (\backslash i \rightarrow \text{map } (+ (i+1)) (\text{iota } i)) \text{ arr}) \equiv$   
1. let ip1s = map (\i -> i+1) arr in -- [2, 3, 4, 5]  
2. let iots =  $\mathcal{F}(\text{map } (\backslash i \rightarrow (\text{iota } i)) \text{ arr})$  in  
3. let ip1rs=  $\mathcal{F}(\text{map2 } (\backslash i \text{ ip1} \rightarrow (\text{replicate } i \text{ ip1})) \text{ arr ip1s})$   
4. in  $\mathcal{F}(\text{map2 } (\backslash \text{ip1r iot} \rightarrow \text{map2 } (+) \text{ ip1r iot}) \text{ ip1rs iots})$ 
```

Applying the new rules results in:

```
1.  $S_{res}^1 = \text{arr}$  -- arr = [1, 2, 3, 4]  
2. ( $B_{res}^1, F'_{res}$ ) = mkFlagArray arr 0 (iota (length arr))  
--  $B_{res}^1 = [0, 1, 3, 6]$   
3.  $F_{res} = \text{map bool.u32 } F'_{res}$   
4.  $II_{res}^1 = \text{sgmScan}^{inc} (+) 0 F_{res} F'_{res}$  -- [0, 1,1, 2,2,2, 3,3,3,3]  
5. ip1s = map (\ i -> i+1 ) arr -- [2, 3, 4, 5]  
6. iots = map2 (\ ind sgm -> ind -  $B_{res}^1[\text{sgm}]$  ) (iota flenres)  $II_{res}^1$   
-- =  $II_{arr}^2 = [0, 0,1, 0,1,2, 0,1,2,3]$   
7. ip1rs= map (\ sgm -> ip1s[sgm] )  $II_{res}^1$  -- [2, 3,3, 4,4,4, 5,5,5,5]  
8. in map2 (+) ip1rs iots -- [2, 3,4, 4,5,6, 5,6,7,8]
```

Lines 6 – 8 are trivially fusable \Rightarrow the iots and ip1rs arrays are not manifested in memory.

Revisiting Our Example: Think Like a Human

I. Normalize the code:

```
map (\i -> let ip1 = i+1 in
        let iot = (iota i) in
        let ip1r = (replicate i ip1)
        in map2 (+) ip1r iot
    ) arr
```

Using the new intuition results in:

```
1.  $S_{res}^1 = arr$  -- arr = [1, 2, 3, 4]
2.  $(B_{res}^1, F'_{res}) = mkFlagArray\ arr\ 0\ (iota\ (length\ arr))$ 
3.  $F_{res} = map\ bool.u32\ F'_{res}$ 
4.  $II_{res}^1 = sgmScan^{inc}\ (+)\ 0\ F_{res}\ F'_{res}$ 
5. in map2 (\ sgm ind -> let ip1      = arr[sgm] + 1
6.                      let iot_el = ind -  $B_{res}^1[sgm]$ 
7.                      in ip1 + iot_el
8.      )  $II_{res}^1\ (iota\ (length\ II_{res}^1))$ 
```

Have done a tiny bit better job than the compiler, as array `ip1s` is not manifested either.

PERFORMANCE DEMONSTRATION

Parallel Basic Blocks Recap

Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening
Several Re-Write Rules (inefficient for replicate & iota)
Jagged (Irregular Multi-Dim) Array Representation
Revisiting the Rewrites for Replicate & Iota Nested Inside Map
Revisiting the Solution to Our Example

Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules

Flattening by Function Lifting
Flattening Quicksort
Flattening Prime-Number (Sieve) Computation

Nested vs Flattened Parallelism: Reduce Inside Map

(5) Reduce Inside a Map or Segmented Reduce:

```
let arr = [[1, 3, 4], [6, 7]] in  
let res = map (\x -> reduce (+) 0 x) arr  
-- should result in [8, 13]
```

Nested vs Flattened Parallelism: Reduce Inside Map

(5) Reduce Inside a Map or Segmented Reduce:

```
let arr = [[1, 3, 4], [6, 7]] in
let res = map (\x -> reduce (+) 0 x) arr
-- should result in [8, 13]
```

translates to a **scan-pack** composition:

1. the length of `res` equals the number of subarrays of `arr`;
2. the shape of `arr` is scanned: the result records the position of the last element in a segment plus one;
3. segmented scan is applied on the input array: the last element in a segment holds the reduced value of the segment;
4. segment's last element is extracted by a map operation.

```

F(res = map (\row -> reduce  $\odot$  0 $\odot$  row) arr) =>
-- Sarr0 = [2], Sarr1 = [3, 2], Farr = [1, 0, 0, 1, 0], Darr = [1, 3, 4, 6, 7]
1. Sres0 = Sarr0 -- Sres0 = [2]
2. indsp1 = scan (+) 0 Sarr1 -- indsp1 = [3, 5]
3. tmp = sgmScan ( $\odot$ ) 0 $\odot$  Farr Darr -- tmp = [1, 4, 8, 6, 13]
4. Dres = map2(\s ip1 -> if s ≤ 0 then 0 $\odot$ 
                    else tmp[ip1-1]) Sarr1 indsp1 -- Dres = [8, 13]

```

Nested vs Flattened Parallelism: Reduce Inside Map

(5) Reduce Inside a Map or Segmented Reduce:

```
let arr = [[1, 3, 4], [6, 7]] in
let res = map (\x -> reduce (+) 0 x) arr
-- should result in [8, 13]
```

Nested vs Flattened Parallelism: Reduce Inside Map

(5) Reduce Inside a Map or Segmented Reduce:

```
let arr = [[1, 3, 4], [6, 7]] in  
let res = map (\x -> reduce (+) 0 x) arr  
-- should result in [8, 13]
```

We can also “cheat” and use a histogram-like computation

Nested vs Flattened Parallelism: Reduce Inside Map

(5) Reduce Inside a Map or Segmented Reduce:

```
let arr = [[1, 3, 4], [6, 7]] in
let res = map (\x -> reduce (+) 0 x) arr
-- should result in [8, 13]
```

We can also “cheat” and use a histogram-like computation

$\mathcal{F}(\text{res} = \text{map } (\backslash \text{row} \rightarrow \text{reduce } \odot 0 \odot \text{row}) \text{ arr}) \Rightarrow$

-- $S_{arr}^0 = [2]$, $S_{arr}^1 = [3, 2]$, $F_{arr} = [1, 0, 0, 1, 0]$, $D_{arr} = [1, 3, 4, 6, 7]$

1. $S_{res}^0 = S_{arr}^0$ -- $S_{res}^0 = [2]$

2. $D_{res} = \text{hist } (\odot) 0 \odot (S_{arr}^0[0]) \parallel_{arr}^1 D_{arr}$

How else can I try to optimize this code by hand?

Nested vs Flattened Parallelism: Reduce Inside Map

(5) Reduce Inside a Map or Segmented Reduce:

```
let arr = [[1, 3, 4], [6, 7]] in
let res = map (\x -> reduce (+) 0 x) arr
-- should result in [8, 13]
```

We can also “cheat” and use a histogram-like computation

$\mathcal{F}(\text{res} = \text{map } (\backslash \text{row} \rightarrow \text{reduce } \odot 0 \odot \text{row}) \text{ arr}) \Rightarrow$
 $-- S_{arr}^0 = [2], S_{arr}^1 = [3, 2], F_{arr} = [1, 0, 0, 1, 0], D_{arr} = [1, 3, 4, 6, 7]$
 1. $S_{res}^0 = S_{arr}^0$ $-- S_{res}^0 = [2]$
 2. $D_{res} = \text{hist } (\odot) 0 \odot (S_{arr}^0[0]) \parallel_{arr}^1 D_{arr}$

How else can I try to optimize this code by hand?

- practical performance refers to how many global-memory accesses you perform
- accessing II_{arr}^1 from memory has significant cost
- in some practical cases, it might be more efficient to not manifest II_{arr}^1 , but instead to compute its elements by binary searching the B_{arr}^1 array.

Flattening Scatter and Histogram

How does one flattens a scatter perfectly nested inside a map?

How does one flattens a histogram perfectly nested inside a map?

Flattening Scatter and Histogram

How does one flattens a scatter perfectly nested inside a map?

How does one flattens a histogram perfectly nested inside a map?

You will have to answer it yourselves as part of the third weekly assignment :)

Treating a Scalar Variant to the Outer Map

(6) The inner construct uses a scalar variant to the outer map:

```
let res = map2 (\x ys -> map (+x) ys) [1,3] [[4,5,6], [9,7]] ≡  
let res = [map (+1) [4,5,6], map (+3) [9,7]]  
let res = [ [5,6,7], [12,10] ]
```

Treating a Scalar Variant to the Outer Map

(6) The inner construct uses a scalar variant to the outer map:

```
let res = map2 (\x ys -> map (+x) ys) [1,3] [[4,5,6], [9,7]] ≡  
let res = [map (+1) [4,5,6], map (+3) [9,7]]  
let res = [ [5,6,7], [12,10] ]
```

Traditionally, this is handled by expanding (replicating) each x across the whole segment

```
let Dxss = [1, 1, 1, 3, 3]  
let res = map2 (+) [1, 1, 1, 3, 3 ]  
                  [4, 5, 6, 9, 7 ]  
                  = = = = =  
                  [5, 6, 7, 12, 10]
```

Instead, we use II_{arr}^1 to indirectly access in the xs array:

```
 $\mathcal{F}(\text{res} = \text{map2 } (\lambda x \text{ ys} \rightarrow \text{map } (f \ x) \text{ ys}) \text{ xs yss}) \Rightarrow$   
-- xs = [1,3],  $S_{yss}^1 = [3,2]$ ,  $F_{yss} = [1,0,0,1,0]$ ,  $D_{yss} = [4,5,6,9,7]$   
1.  $S_{res}^1 = S_{yss}^1$   
2.  $D_{res} = \text{map2 } (\lambda y \text{ sgmind} \rightarrow f \text{ xs[sgmind] } y) \text{ } D_{yss} \text{ II}_{yss}^1$   
--  $\text{II}_{yss}^1 = [0,0,0,1,1]$ ,  $D_{res} = [5,6,7,12,10]$ 
```

Treating Indexing Variant to the Outer Map

(7) Indexing Operations Variant to the Outer Map:

```
let res = map2 (\i xs -> xs[i]) [2,0] [[4,5,6], [9,7]] ≡  
let res = [ 6, 9 ]
```

Treating Indexing Variant to the Outer Map

(7) Indexing Operations Variant to the Outer Map:

```
let res = map2 (\i xs -> xs[i]) [2,0] [[4,5,6], [9,7]] ≡  
let res = [ 6, 9 ]
```

To corresponding flat index in $D_{y_{ss}}$ is obtained by summing up

- the start offset of every segment, which we get from $B_{y_{ss}}^1$, and
- the index inside the segment, which we get from i_s

$\mathcal{F}(\text{res} = \text{map2 } (\lambda i \text{ xs} \rightarrow \text{xs}[i]) \text{ is } x_{ss}) \Rightarrow$

-- $i_s = [2,0]$, $S_{x_{ss}}^1 = [3,2]$, $B_{x_{ss}}^1 = [0,3]$, $D_{x_{ss}} = [4,5,6,9,7]$

1. $S_{res}^0 = S_{i_s}^0 \text{ -- } = S_{i_s}^0 = [2]$

2. $D_{res} = \text{map2 } (\lambda \text{ off } i \rightarrow D_{x_{ss}}[\text{off} + i]) B_{x_{ss}}^1 \text{ is -- } D_{res} = [6, 9]$

Nested vs Flattened Parallelism: If Inside a Map 2D Case

(8) If-Then-Else with inner parallelism nested inside a map:

```
bs  = [F,T,F,T]
xss = [[1,2,3],[4,5,6,7],[8,9],[10]]
res = map(\b xs -> if b then map (+1) xs else map (*2) xs) bs xss
res = [ map(*2)[1,2,3], map(+1)[4,5,6,7], map(*2)[8,9], map(+1)[10] ]
res = [ [2,4,6], [5,6,7,8], [16,18], [11] ]
```

Nested vs Flattened Parallelism: If Inside a Map 2D Case

(8) If-Then-Else with inner parallelism nested inside a map:

```
bs = [F,T,F,T]
xss = [[1,2,3],[4,5,6,7],[8,9],[10]]
res = map(\b xs -> if b then map (+1) xs else map (*2) xs) bs xss
res = [ map(*2)[1,2,3], map(+1)[4,5,6,7], map(*2)[8,9], map(+1)[10] ]
res = [ [2,4,6], [5,6,7,8], [16,18], [11] ]
```

translates to a **scatter-map-gather** composition. Intuition:

1. compute `iinds`, the permutation of segments w.r.t. `bs`;
 - 2-3. partition the `xss` array based on `bs`;
 - 4-5. **flatten outer map and/on top of the parallel code of the then and else branches**;
 6. inverse permute the resulted segments according to `iinds`.
- ```
1. iinds = partition2 (\i -> bs[i]) (iota (length b)) -- [1,3,0,2]
2. xssthen = gatherThen iinds xss -- ([4,1], [4,5,6,7, 10])
3. xsselse = gatherElse iinds xss -- ([3,2], [1,2,3, 8,9])
-- Recursively Flatten the Then and Else Branches!
4. resthen = F(map (map (+1)) xssthen) -- ([4,1], [5,6,7,8, 11])
5. reselse = F(map (map (*2)) xsselse) -- ([3,2], [2,4,6,16,18])
6. res = inversePermute iinds (resthen++reselse)
-- ([3,4,2,1], [2,4,6, 5,6,7,8, 16,18, 11])
```

# Nested vs Flattened Parallelism: If Inside a Map 2D Case

## (8) If-Then-Else with inner parallelism nested inside a map:

```
bs = [F,T,F,T], xss = [[1,2,3],[4,5,6,7],[8,9],[10]], S1xss=[3,4,2,1], f=map (+1), g=map (*2)

F(res = map2 (\b xs => if b then f xs else g xs) bs xss) =>
(spl, iinds) = partition2 bs (iota (length bs)) -- (2, [1,3,0,2])
(S1xssthen, S1xsselse) = split spl (map (\ii => S1xss[ii]) iinds) -- ([4,1],[3,2])
maskxss = map (\sgmind => bs[sgmind]) ||1xss -- [F,F,F,T,T,T,T,F,F,T]
(brk, Dpxss) = partition2 maskxss Dxss
(Dxssthen, Dxsselse) = split brk Dpxss -- ([4,5,6,7,10],[1,2,3,8,9])
(S1resthen, Dresthen) = F(map f) (S1xssthen, Dxssthen) -- ([4,1], [5,6,7,8,11])
(S1reselse, Dreselse) = F(map g) (S1xsselse, Dxsselse) -- ([3,2], [2,4,6,16,18])
S1Pres = S1resthen ++ S1reselse -- [4,1,3,2]
S1res = scatter (replicate (length bs) 0) iinds S1Pres -- [3,4,2,1]
B1res = scanexc (+) 0 S1res -- [0,3,7,9]
FPres = mkFlagArray S1Pres 0 (map (+1) iinds) -- [2,0,0,0,4,1,0,0,3,0]
||1Pres = sgmscan (+) 0 FPres FPres |> map (\x => x-1) -- [1,1,1,1,3,0,0,0,2,2]
||2Pres = ||2resthen ++ ||2reselse -- [0,1,2,3,0, 0,1,2,0,1]
sindsres = map2 (\sgm iin => B1res[sgm] + iin) ||1Pres ||2Pres
-- [3+0,3+1,3+2,3+3, 9+0, 0+0,0+1,0+2, 7+0,7+1]=[3,4,5,6,9,0,1,2,7,8]
Dres = scatter (replicate flenres 0) sindsres (Dresthen ++ Dreselse)
-- [2,4,6, 5,6,7,8, 16,18, 11]
(S1res, Dres)
```

# Nested vs Flattened Parallelism: Do Loop Inside a Map

## (9) Flattening a Do Loop Nested Inside a Map:

- compute the maximal loop count  $n_{max}$
- interchange the loop and the map:
  - ▶ loop count becomes  $n_{max}$
  - ▶ the loop body is wrapped inside a `if i < n` condition, and
  - ▶ the new loop body is flattened!

$\mathcal{F}(\text{res} = \text{map2} (\backslash n \text{ xs} \rightarrow \text{loop} (\text{xs}) \text{ for } i < n \text{ do } f \text{ xs}) \text{ ns } \text{xss}) \Rightarrow$

1.  $n_{max} = \text{reduce } \max \text{ } 0 \text{ } i32 \text{ ns}$
2.  $g \text{ i m arr} = \text{if } i < m \text{ then } f \text{ arr } \text{else } \text{arr}$
3.  $\text{loop}(S_{xss}^1, D_{xss}) \text{ for } i < n_{max} \text{ do}$
4.  $\mathcal{F}(\text{map2 } (g \text{ i})) \text{ ns } (S_{xss}^1, D_{xss})$
5.  $-- (g \text{ i})^L \text{ ns } (S_{xss}^1, D_{xss})$

But this treatment does not necessarily preserve the work asymptotic ... what to do?

# Nested vs Flattened Parallelism: Do Loop Inside a Map

## (9) Flattening a Do Loop Nested Inside a Map:

- compute the maximal loop count  $n_{max}$
- interchange the loop and the map:
  - ▶ loop count becomes  $n_{max}$
  - ▶ the loop body is wrapped inside a `if i < n` condition, and
  - ▶ the new loop body is flattened!

$\mathcal{F}(\text{res} = \text{map2} (\backslash n \text{ xs} \rightarrow \text{loop}(\text{xs}) \text{ for } i < n \text{ do } f \text{ xs}) \text{ ns } \text{xss}) \Rightarrow$

1.  $n_{max} = \text{reduce } \max \text{ } 0 \text{ i } 32 \text{ ns}$
2.  $g \text{ i } m \text{ arr} = \text{if } i < m \text{ then } f \text{ arr } \text{else } \text{arr}$
3. **loop** ( $S_{xss}^1$ ,  $D_{xss}$ ) **for**  $i < n_{max}$  **do**
4.      $\mathcal{F}(\text{map2} (g \text{ i})) \text{ ns } (S_{xss}^1, D_{xss})$
5.     --  $(g \text{ i})^L \text{ ns } (S_{xss}^1, D_{xss})$

But this treatment does not necessarily preserve the work asymptotic ... what to do?

If the size of the result can be deduced/inferred:

- allocate the flat-result array before the loop
- filter out the empty segments, and make the loop iterate until the shape is empty
- each time a segment finish execution (1) it is scattered into the result, and (2) it is filtered out from the running set of segments.

## Parallel Basic Blocks Recap

### Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening  
Several Re-Write Rules (inefficient for replicate & iota)  
Jagged (Irregular Multi-Dim) Array Representation  
Revisiting the Rewrites for Replicate & Iota Nested Inside Map  
Revisiting the Solution to Our Example

### Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules  
**Flattening by Function Lifting**  
Flattening Quicksort  
Flattening Prime-Number (Sieve) Computation

## Flattening by Function Lifting: Basic Idea

Assume a simple function  $f$ :

```
let f (x: i32) : i32 = x + 1
```

$f$  lifted, denoted  $f^L$  semantically corresponds to `map f`, where the arguments have been expanded to an extra array dimension, and the inner operators/functions have also been lifted:

```
let $+^L$ [n] (as: [n]i32) (bs: [n]i32) : [n]i32 =
 map2 (+) as bs
```

```
let f^L [n] (xs: [n]i32) : [n]i32 =
 xs $+^L$ (replicate n 1)
```

# Flattening by Function Lifting: Basic Idea

Assume a simple function  $f$ :

```
let f (x: i32) : i32 = x + 1
```

$f$  lifted, denoted  $f^L$  semantically corresponds to `map f`, where the arguments have been expanded to an extra array dimension, and the inner operators/functions have also been lifted:

```
let +L [n] (as: [n]i32) (bs: [n]i32) : [n]i32 =
 map2 (+) as bs
```

```
let fL [n] (xs: [n]i32) : [n]i32 =
 xs +L (replicate n 1)
```

- Locals such as  $x \Rightarrow$  left alone
- Global such as  $+$   $\Rightarrow$  lifted  $(+^L)$
- Constants such as  $k \Rightarrow \text{replicate (length xs) } k$ 
  - ▶ good for vectorization, bad for locality, asymptotics
  - ▶ for GPU better to indirectly index into a smaller array, rather than `replicate`.



## Flattening by Function Lifting: Key Insight!

```
let f (xs: []i32) : []i32 = map g xs -- = g^L xs
let f^L (xss: [][]i32) : [][]i32 = (g^L)L -- ???
```

How do we stop lifting?  $g$  and  $g^L$  are enough: no need for  $(g^L)^L$ !

## Flattening by Function Lifting: Key Insight!

```
let f (xs: []i32) : []i32 = map g xs -- = g^L xs
let f^L (xss: [][]i32) : [][]i32 = (g^L)L -- ???
```

How do we stop lifting?  $g$  and  $g^L$  are enough: no need for  $(g^L)^L$ !

```
let f (xs: []i32) : []i32 = map g xs -- = g^L xs
-- in nested parallel form
let f^L (xss: [][]i32) : [][]i32 =
 segment xss (g^L (concat xss))
-- in flatten form
let f^L (S_{xss}^1 : []i32, D_{xss} : []i32) : ([]i32, []i32) =
 (S_{xss}^1 , g^L D_{xss})
```

In Haskell Notation:

|         |    |       |    |           |             |
|---------|----|-------|----|-----------|-------------|
| concat  | :: | [[a]] | -> | [a]       |             |
| segment | :: | [[a]] | -> | [b]       | -> [[b]]    |
|         |    | shape |    | flat data | nested data |

## Parallel Basic Blocks Recap

### Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening  
Several Re-Write Rules (inefficient for replicate & iota)  
Jagged (Irregular Multi-Dim) Array Representation  
Revisiting the Rewrites for Replicate & Iota Nested Inside Map  
Revisiting the Solution to Our Example

### Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules  
Flattening by Function Lifting  
**Flattening Quicksort**  
Flattening Prime-Number (Sieve) Computation

# Recounting Quicksort

Recount the classic nested-parallel definition:

```
let quicksort [n] (arr : [n]f32) : [n]f32 =
 if n < 2 then arr else
 let i = getRand (0, (length arr) - 1)
 let a = arr[i]
 let s1 = filter (< a) arr
 let s2 = filter (== a) arr
 let s3 = filter (> a) arr
 in (quicksort s1) ++ s2 ++ (quicksort s3)
 -- can be re-written as:
 -- rs = map nestedQuicksort [s1, s3]
 -- in (rs[0]) ++ s2 ++ (rs[1])
```

Note: Futhark does not support recursive calls, hence not valid code!

# Nested-Parallel Quicksort Simplified

For simplicity we will rewrite it in terms of `partition2`:

```
let isSorted [n] (as: [n]f32) : bool =
 map (\i -> if i==0 then true else as[i-1] < as[i]) (iota n)
 |> reduce (&&) true

let quicksort [n] (arr: [n]f32) : [n]f32 =
 if isSorted arr then arr else
 let i = getRand (0, (length arr) - 1)
 let a = arr[i]
 let bs = map (< a) arr
 let (q, arr') = partition2 bs 0.0f32 arr
 let (arr<, arr≥) = split q arr'
 in concat <| map quicksort [arr<, arr≥]
```

Note: Futhark does not support recursive calls, irregular map operation, or concat!

## Partition2

Reorders the elements of an array such that those that correspond to a true mask come before those corresponding to false.

```
let partition2 [n] 't (conds: [n] bool) (dummy: t) (arr: [n]t)
 : (i32, [n]t) =
 let tflgs = map (\ c -> if c then 1 else 0) conds
 let fflgs = map (\ b -> 1 - b) tflgs

 let indsT = scan (+) 0 tflgs
 let tmp = scan (+) 0 fflgs
 let lst = if n > 0 then indsT[n-1] else -1
 let indsF = map (+lst) tmp

 let inds = map3 (\ c indT indF -> if c then indT-1 else indF-1)
 conds indsT indsF

 let fltarr = scatter (replicate n dummy) inds arr
 in (lst, fltarr)
```

For example:

```
conds = [F,T,F,T,F,F,T]
xss = [1,2,3,4,5,6,7]
partition2 conds 0 xss => (3, [2,4,7,1,3,5,6])
```

# Lifting Quicksort

**Key Idea: write a function with the semantics of**

`map nestedQuicksort`, i.e., it operates on array of arrays.

```
let isSorted [n] (as: [n]f32) : bool =
 map (\i -> if i==0 then true else as[i-1] < as[i]) (iota n)
 |> reduce (&&) true

let quicksortL (xss: [[]]f32) : [[]]f32 =
 map (\xs ->
 if isSorted xs then xs else -- oh, boy, not if-then-else!
 let i = getRand (0, (length xs) - 1)
 let a = xs[i]
 let bs = map (< a) xs
 let (q, xsp) = partition2 bs 0.0f32 xs
 let (xs<, xs≥) = split q xsp
 in concat <| map quicksort [xs<, xs≥]
) xss
```

Important observations:

- `map quicksort`  $\equiv$  `quicksortL`
- `map(map quicksort)`  $\equiv$  `quicksortL`  $\equiv$  `segment o quicksortL o concat`
- the flat data of `[xs<, xs≥]`  $\equiv$  `xsp`, the result of `partition2`

# Lifting Quicksort

Let us treat the last three lines from the previous implem.:

```
let quicksortL (Sxss1:[] i32 , Dxss:[] f32): ([] i32 , [] f32) = --(xss: [] [] f32)
 if isSorted Dxss then (Sxss1:[] i32 , Dxss:[] f32) else -- big cheat!
 let (Sbss1 , Dbss) = \mathcal{F} (
 map (\xs ->
 let i = getRand (0 , (length xs) - 1)
 let a = xs[i]
 let bs = map (< a) xs
 in bs
) xss
)
 let (ps , (Sxssp1 , Dxssp)) = partition2L Dbss 0.0f32 (Sxss1 , Dxss)
 -- Invariant: Sxssp1 == Sbss1 == Sxss1
 let S[xss<,xss≥]1 = filter (!=0) <| flatten <|
 map2 (λ p s -> if s==0 then [0,0] else [p,s-p]) ps Sxss1
 in quicksortL (S[xss<,xss≥]1 , Dxssp)
```

- S<sub>[xss<a,xss≥a]</sub><sup>1</sup> is the shape of [xs< , xs≥]
- (concat <| quicksort<sup>L</sup>)<sup>L</sup> xsss ≡ concat <| segment xsss <| quicksort<sup>L</sup> (concat xsss) ≡ quicksort<sup>L</sup> (concat xsss)
- The function looks tail recursive now: let's replace it with a loop!



# Lifting Quicksort: Final Implementation

```
let quicksortL [m][n] (S1XSS:[m]i32, DXSS:[n]f32): [n]f32 =
 let (stop, count) = (isSorted DXSS, 0i32)
 let (_, res, _, _) =
 loop (S1XSS, DXSS, stop, count) while (!stop) do
 -- compute helper-representation structures
 let B1XSS = scanexc (+) 0 S1XSS
 let F1XSS = mkFlagArray S1XSS 0i32 <| map (+1) <| iota m
 let ll1XSS = sgmscan (+) 0 F1XSS <|
 map (\f -> if f==0 then 0 else f-1) F1XSS
 -- flattening quicksort:
 let rL = map (\u -> randomInd (0,u-1) count) S1XSS
 let aL = map3 (\r l i -> if l <= 0 then 0.0 else DXSS[B1XSS[i]+r]
) rL S1XSS (iota m)
 let Dbss = map2 (\x sgmind -> aL[sgmind] > x) DXSS ll1XSS
 let (ps, (S1XSSp, DperXSS)) = partition2L Dbss 0.0f32 (S1XSS, DXSS)
 let S1[XSS<,XSS≥] = filter (!=0) <| flatten <|
 map2 (\p s -> if s==0 then [0,0] else [p,s-p]) ps S1XSS
 in (S1[XSS<,XSS≥], DperXSS, isSorted DperXSS, count+1)
 in res
```

PFP Weekly 2 Exercise: Implement partition2<sup>L</sup>

## Parallel Basic Blocks Recap

### Part I: Flattening Nested and Irregular Parallelism

What is "Flattening"? Recipe for Applying Flattening  
Several Re-Write Rules (inefficient for replicate & iota)  
Jagged (Irregular Multi-Dim) Array Representation  
Revisiting the Rewrites for Replicate & Iota Nested Inside Map  
Revisiting the Solution to Our Example

### Part II: Flattening Nested and Irregular Parallelism

More Flattening Rules  
Flattening by Function Lifting  
Flattening Quicksort  
Flattening Prime-Number (Sieve) Computation

# How Does One Flatten Prime Numbers?

## The important bit with nested parallelism:

```
sqrt_primes = primesOpt (sqrt (fromIntegral n))
nested = map (\p -> let m = (n `div` p)
 in map (\j -> j*p) [2..m]
) sqrt_primes
not_primes = reduce (++) [] nested
```

# How Does One Flatten Prime Numbers?

## The important bit with nested parallelism:

```
sqrt_primes = primesOpt (sqrt (fromIntegral n))
nested = map (\p -> let m = (n `div` p)
 in map (\j -> j*p) [2..m]
) sqrt_primes
not_primes = reduce (++) [] nested
```

## Normalize the nested map:

```
sqrt_primes = primesOpt (sqrt (fromIntegral n))
nested = map (\p ->
 let m = n `div` p in -- distribute map
 let mm1 = m - 1 in -- distribute map
 let iot = iota mm1 in -- \mathcal{F} rule 4
 let twom = map (+2) iot in -- \mathcal{F} rule 2
 let rp = replicate mm1 p in -- \mathcal{F} rule 3
 in map (\(j,p) -> j*p) (zip twom rp) -- \mathcal{F} rule 2
) sqrt_primes
not_primes = reduce (++) [] nested -- ignore, already flat
```

Flattening PrimeOpt was part of PMPH's Weekly Assignment 2!